

RAG(Retrieval Augmented Generation)Cheatsheet

Stages in RAG:

- 1. Loading:
 - Import your data (text files, PDFs, databases, APIs) using LlamaHub's extensive range of connectors.
- 2. Indexing:
 - Create searchable data structures, primarily through vector embeddings and metadata strategies, enabling efficient context retrieval.
- 3. Storing:
 - Securely store your indexed data and metadata for quick access without the need to re-index.
- 4. Querying:
 - Utilize LLMs and LlamaIndex data structures for diverse querying techniques, including sub-queries and hybrid strategies.
- 5. Evaluation:
 - Continuously assess the effectiveness of your pipeline to ensure accuracy, faithfulness, and response speed.

Application Types:

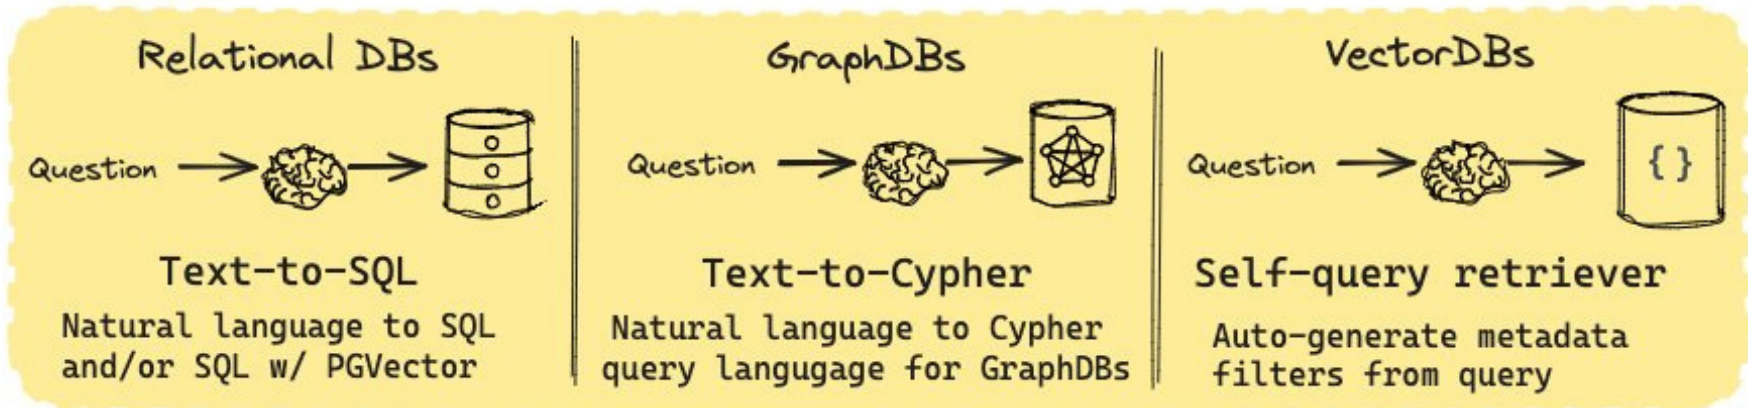
- 1. Query Engines:
 - For direct question-answering over your data.
- 2. Chat Engines:
 - Enables conversations with your data for an interactive experience.
- 3. Agents:
 - Automated decision-makers that interact with external tools, adaptable for complex tasks.

Key Concepts:

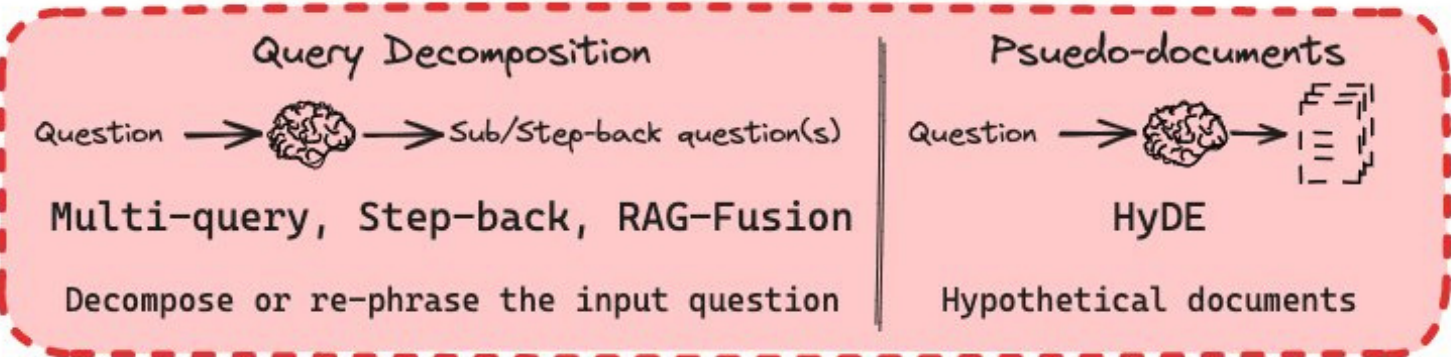
- 1. Nodes and Documents:
 - Fundamental units in LlamaIndex, where Documents encapsulate data sources and Nodes represent data "chunks" with associated metadata.
- 1. Connectors:
 - Bridge various data sources into the RAG framework, transforming them into Nodes and Documents.
- 1. Indexes:
 - The backbone of RAG, enabling the storage of vector embeddings in a vector store along with crucial metadata.
- 1. Embeddings:
 - Numerical representations of data, facilitating the relevance filtering process.
- 1. Retrievers:
 - Define efficient retrieval strategies, ensuring the relevancy and efficiency of data retrieval.
- 1. Routers:
 - Manage the selection of appropriate retrievers based on query specifics and metadata.
- 1. Node Postprocessors:
 - Apply transformations or re-ranking logic to refine the set of retrieved nodes.
- 1. Response Synthesizers:
 - Craft responses from the LLM, utilizing user queries and retrieved text chunks for enriched answers.



Query Construction



Query Translation



Routing

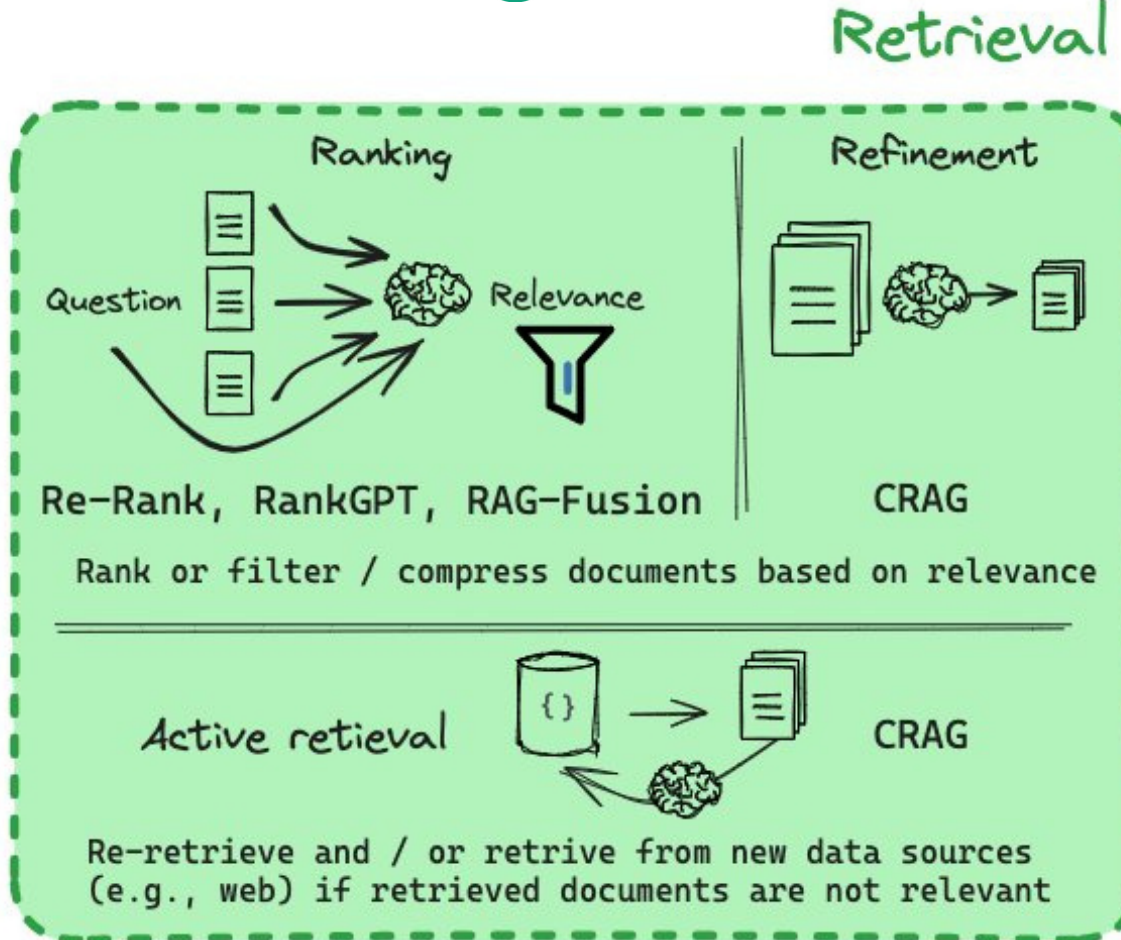
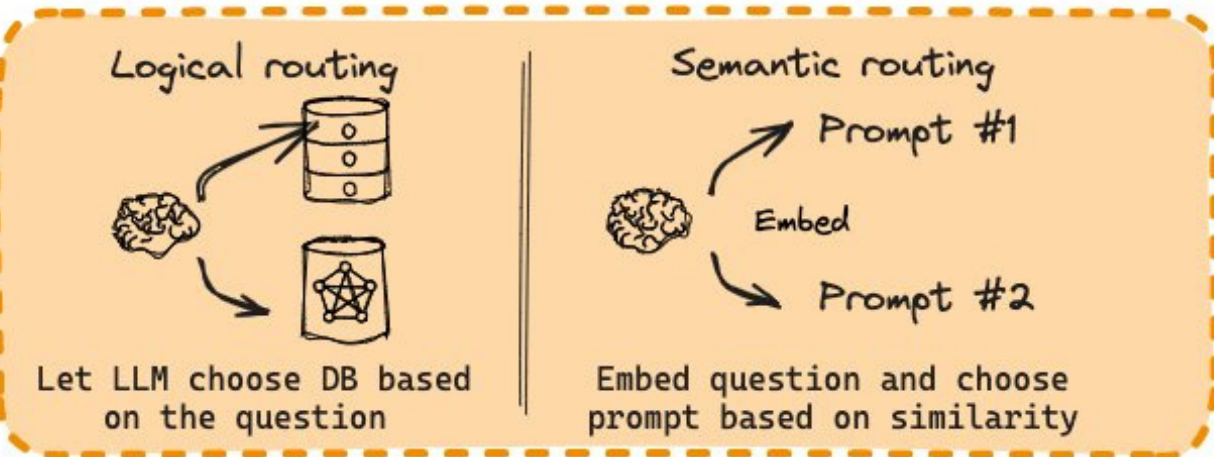
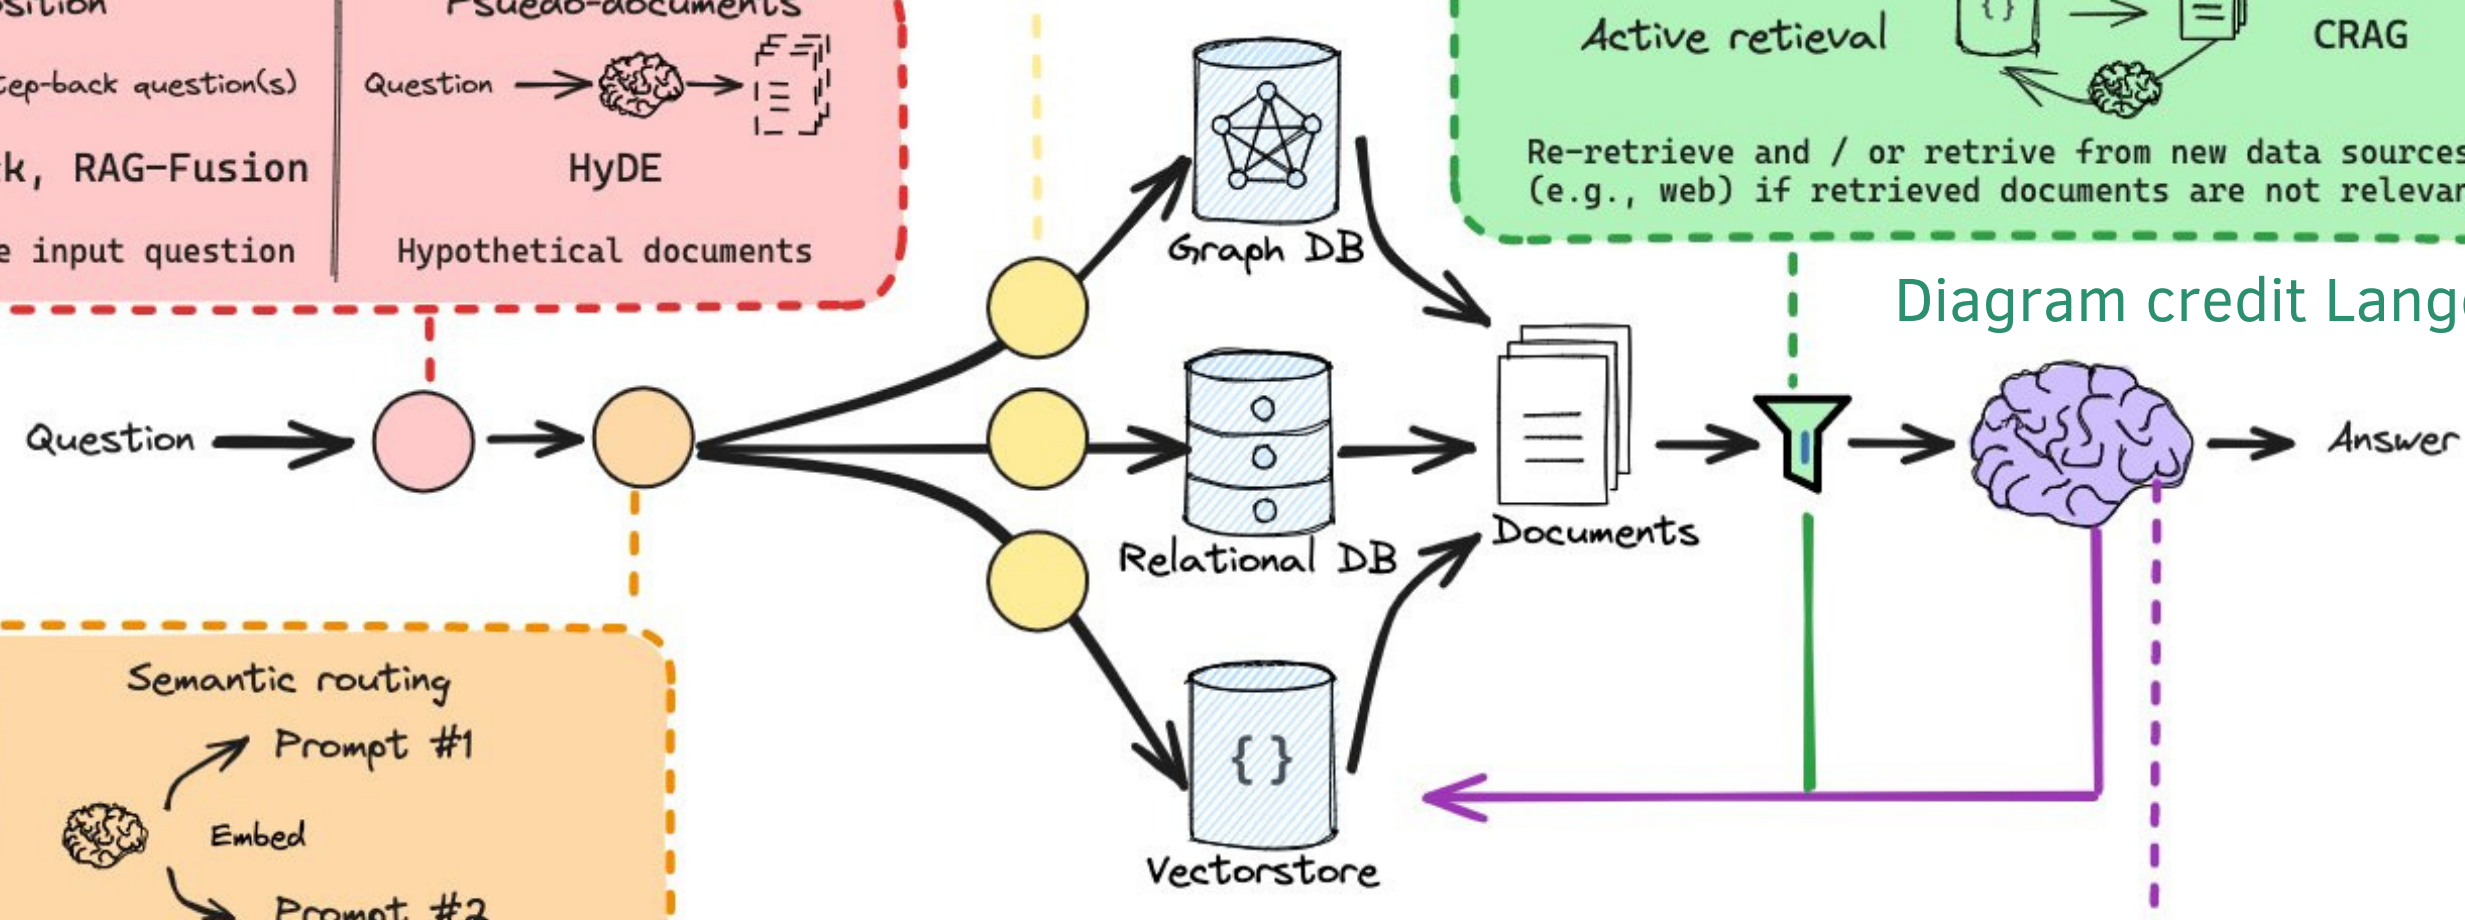
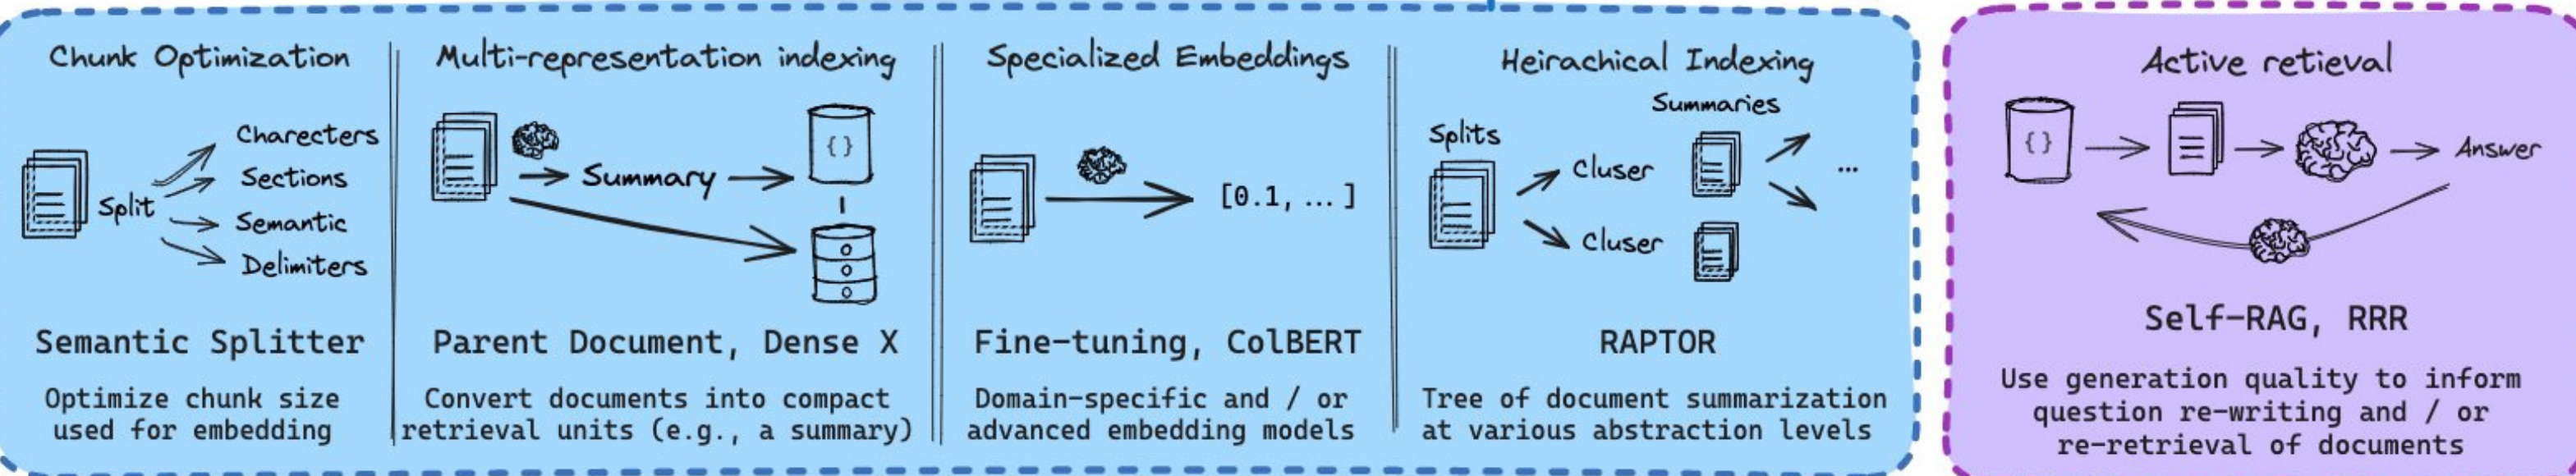


Diagram credit Langchain



Indexing

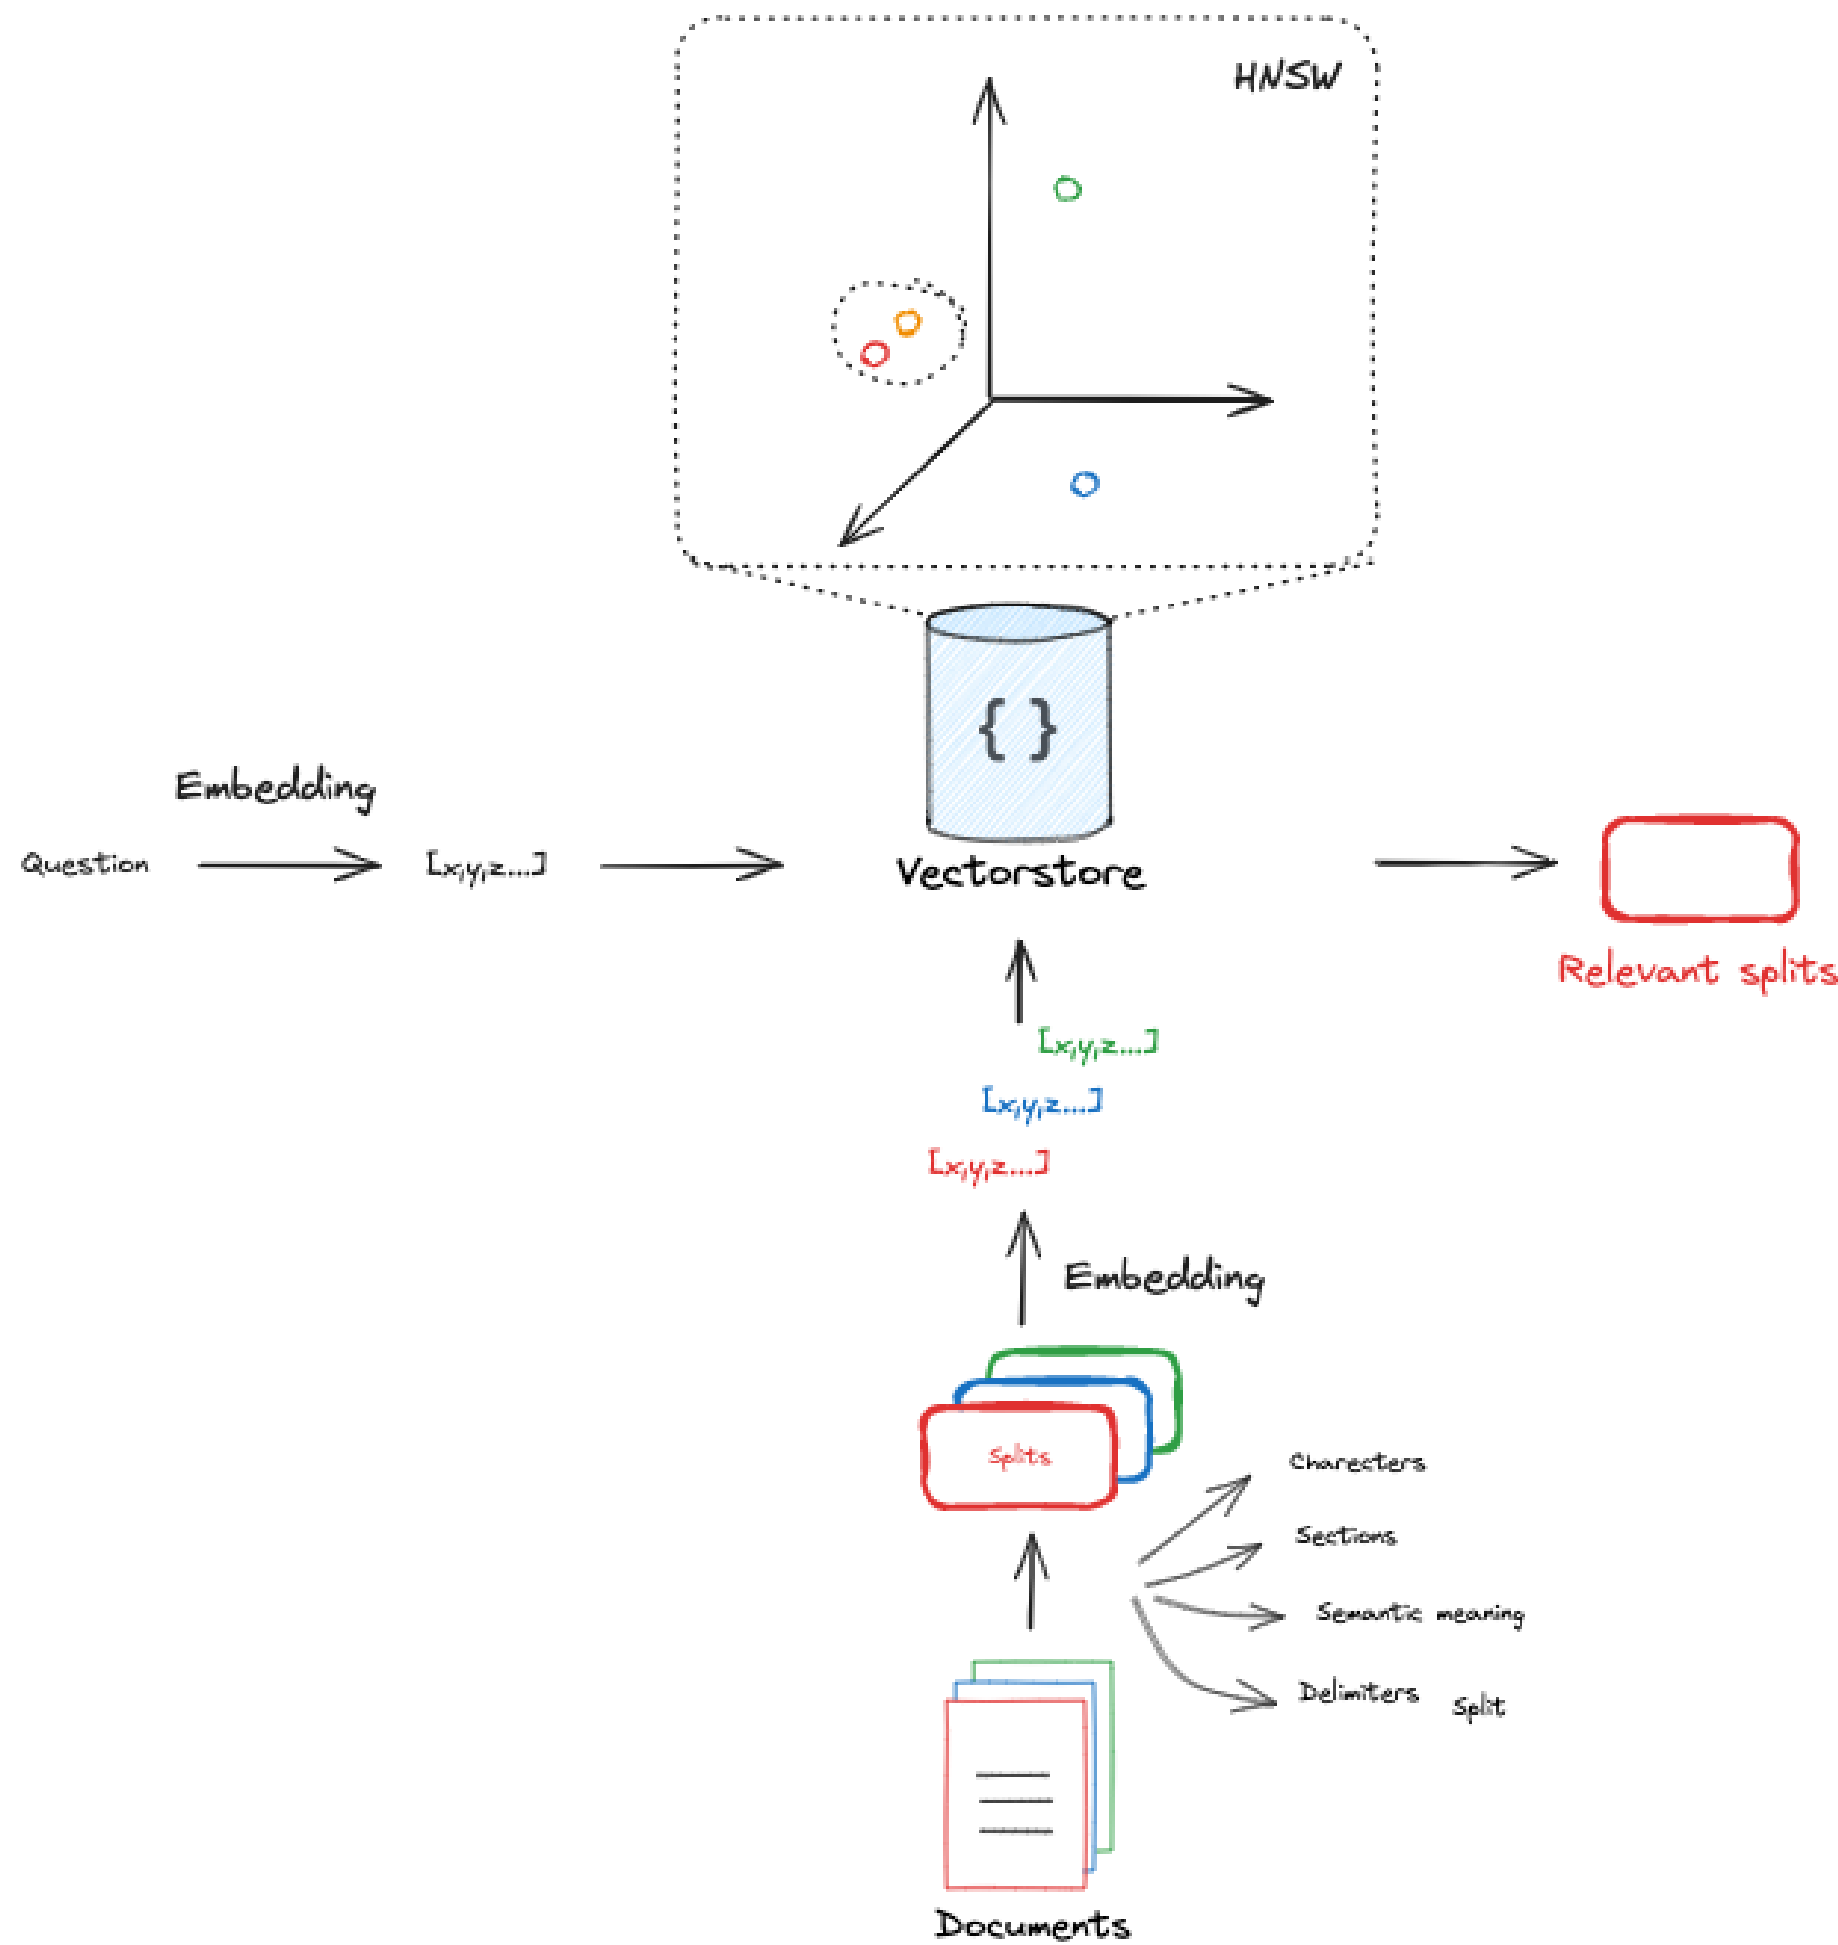
Generation



Steve Nouri

RAG (Retrieval Augmented Generation) Cheatsheet

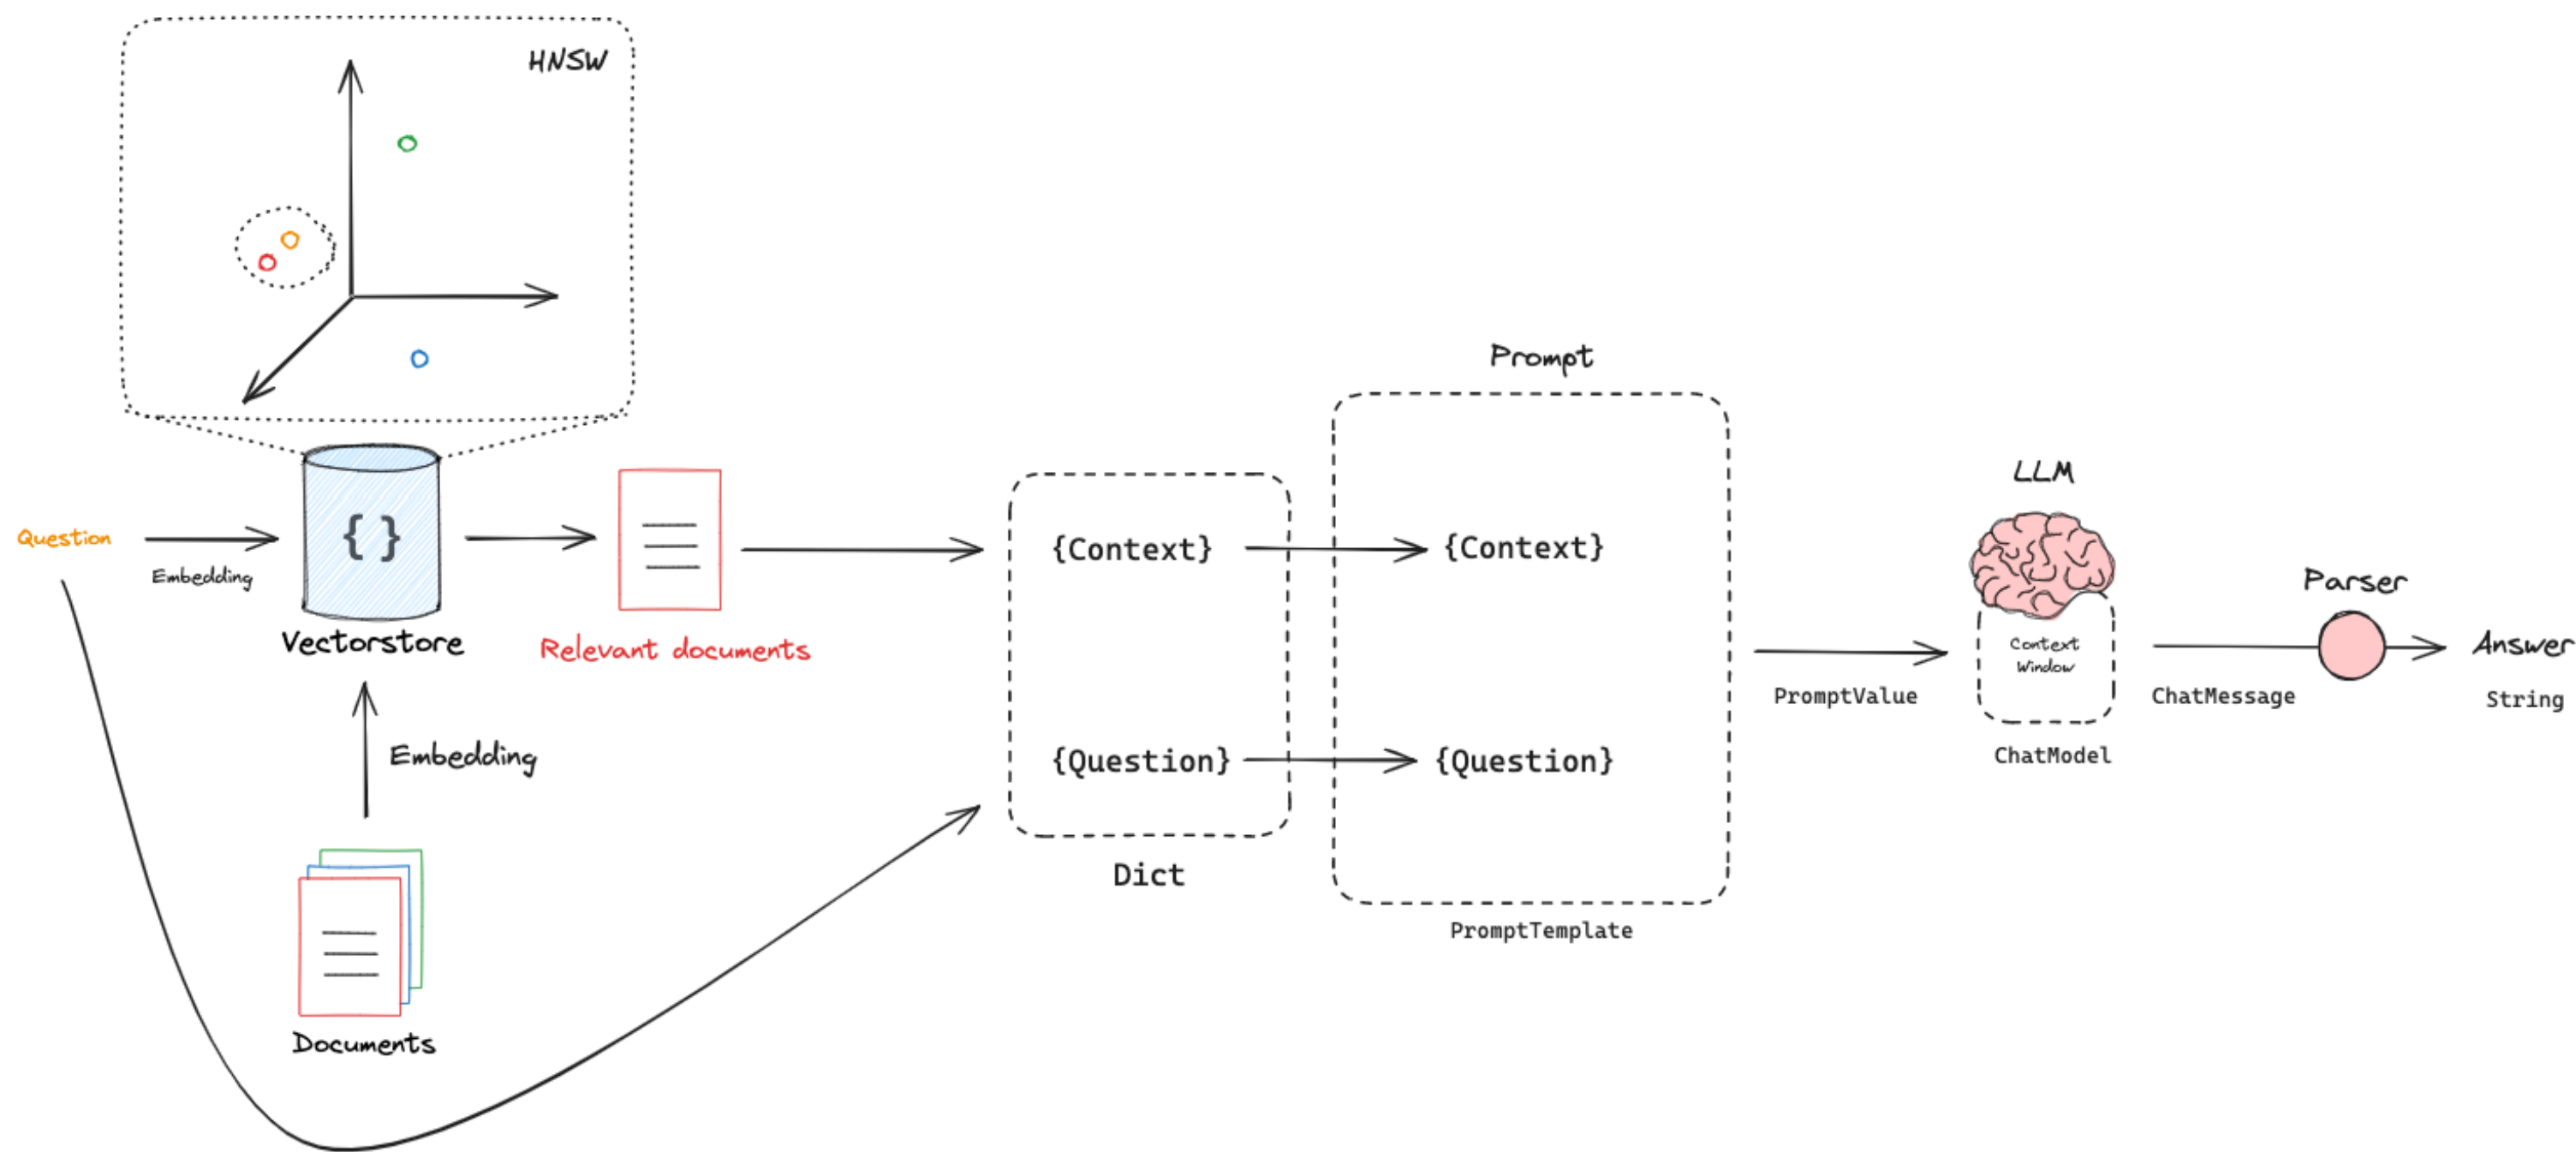
Indexing:



Steve Nouri

https://github.com/langchain-ai/rag-from-scratch/blob/main/rag_from_scratch_1_to_4.ipynb

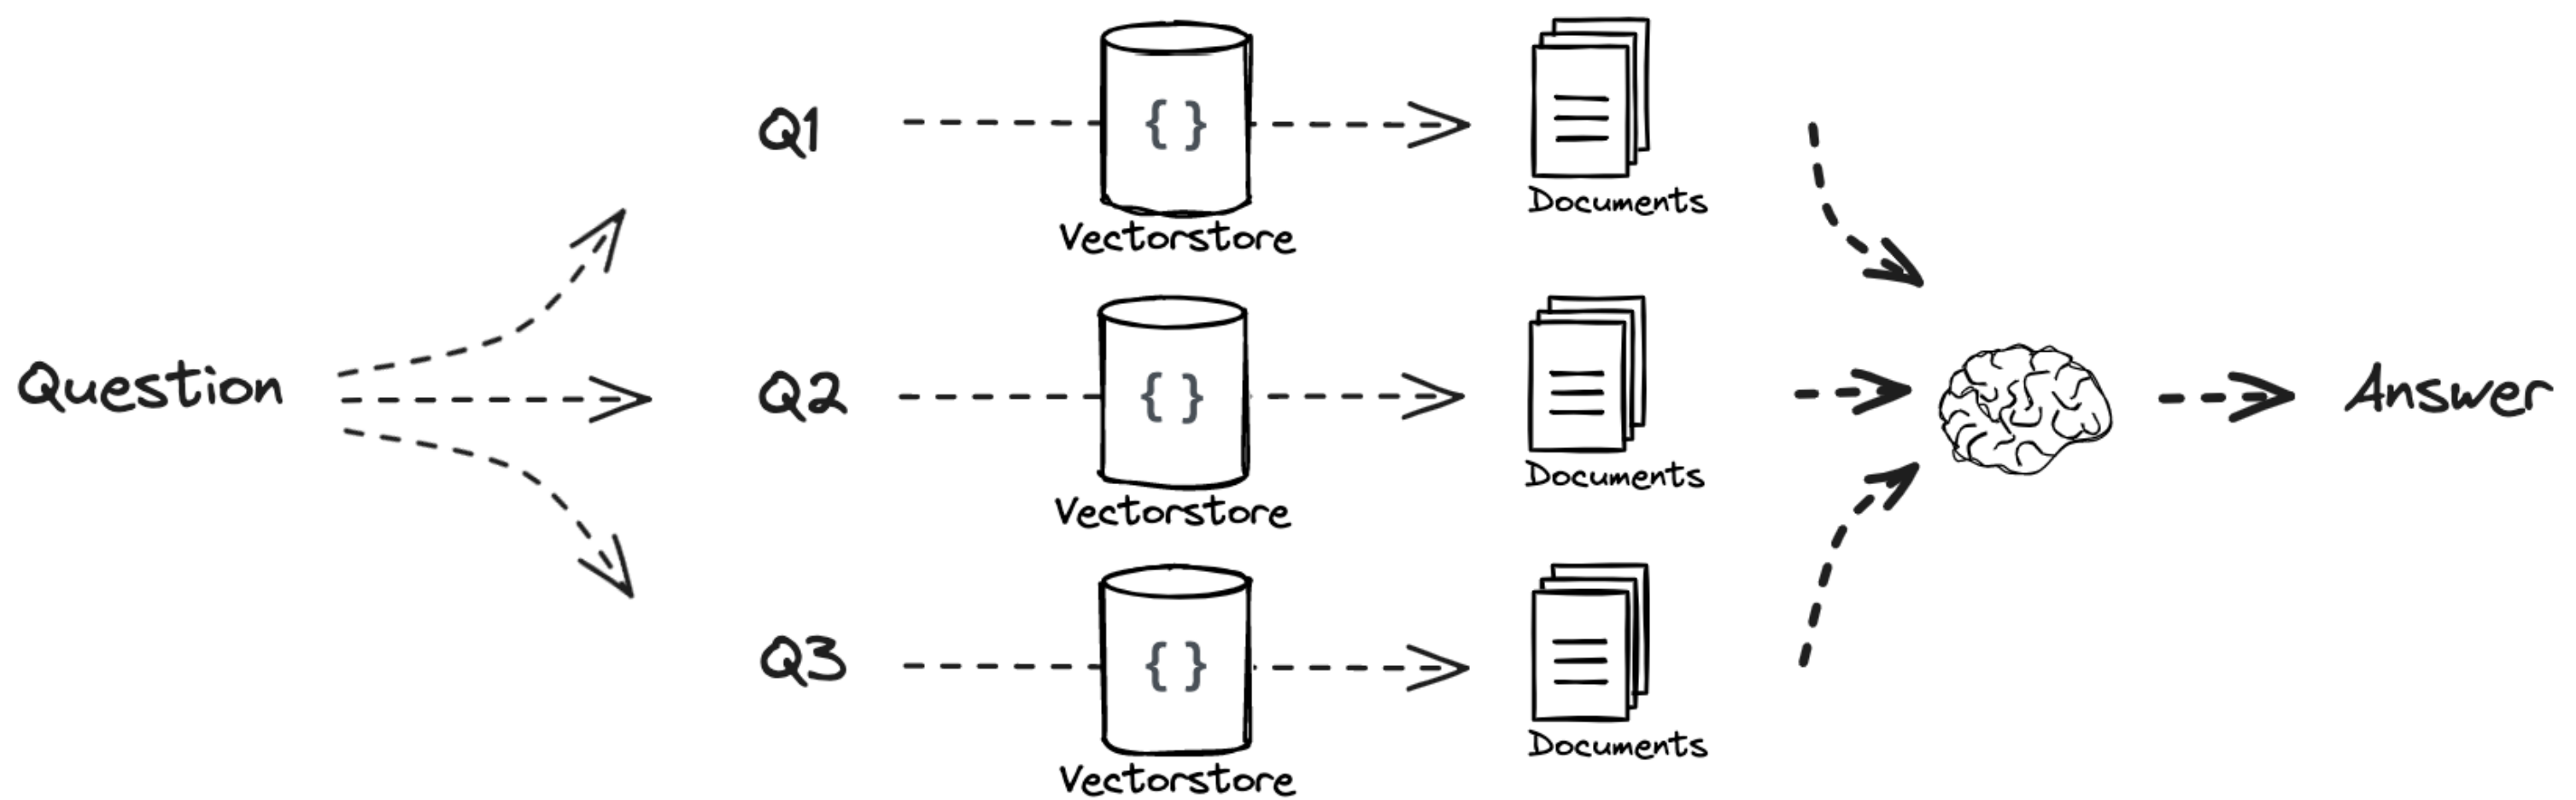
Generation:



https://github.com/langchain-ai/rag-from-scratch/blob/main/rag_from_scratch_1_to_4.ipynb

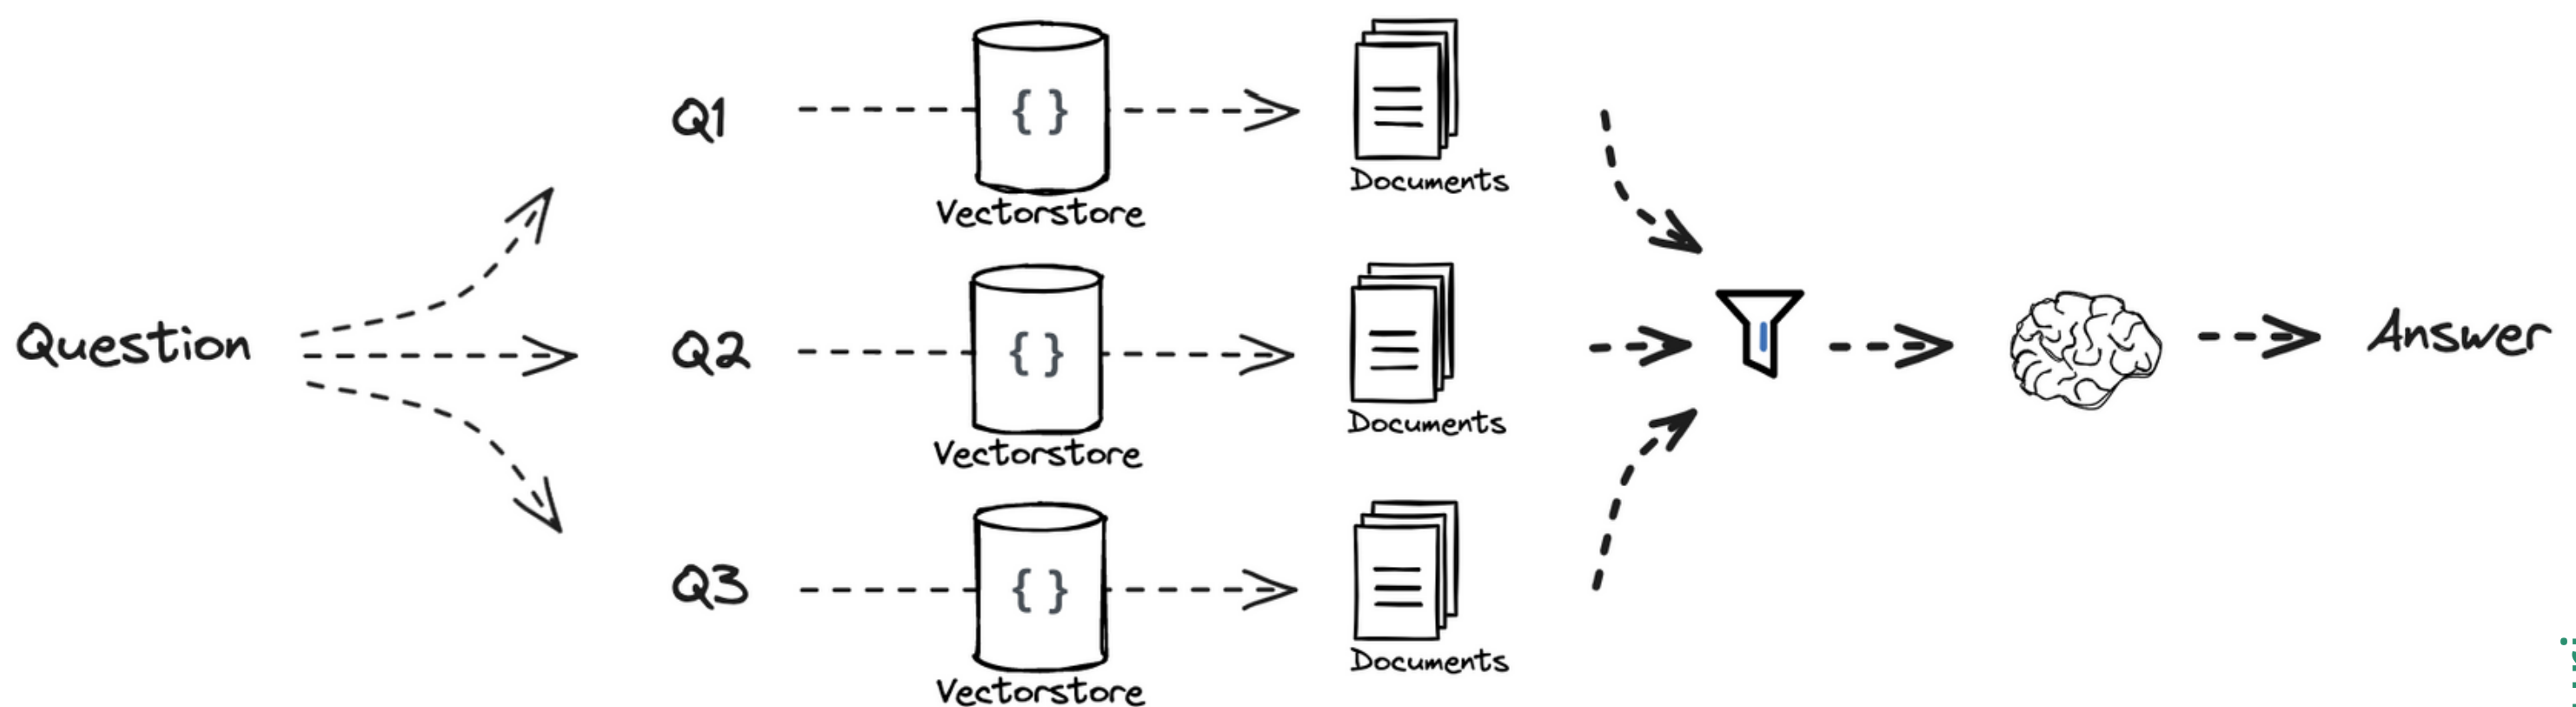
RAG (Retrieval Augmented Generation) Cheatsheet

Multi Query:



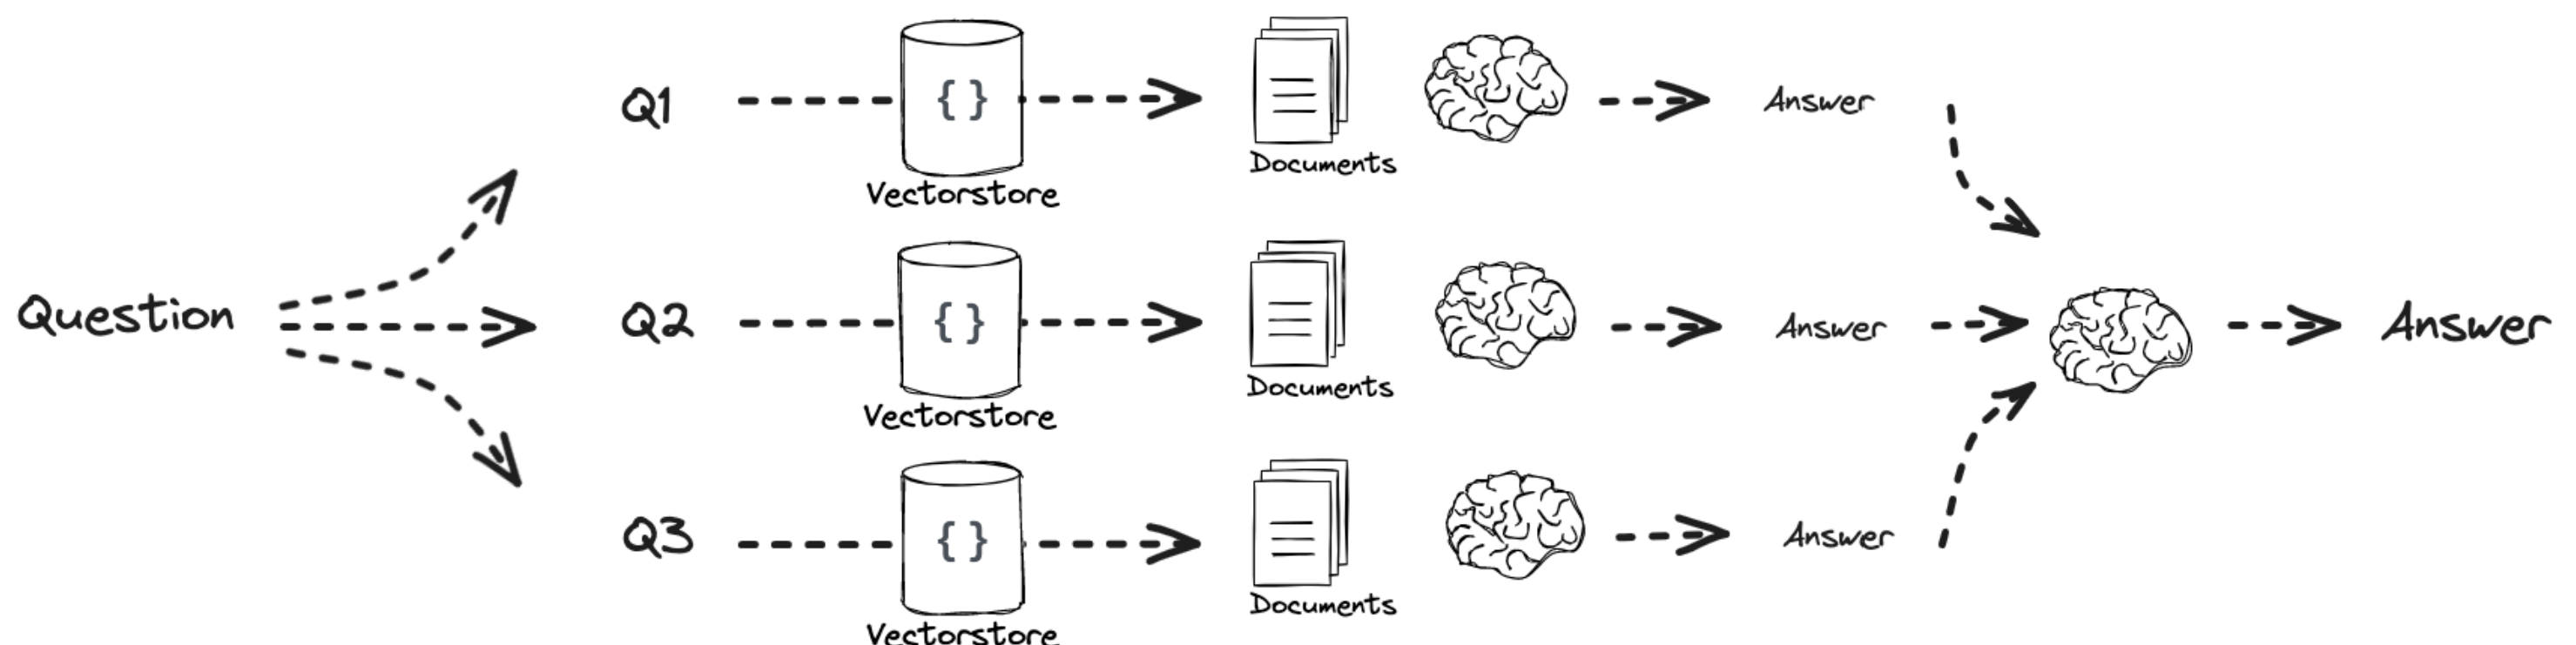
https://python.langchain.com/docs/modules/data_connection/retrievers/MultiQueryRetriever

RAG-Fusion:



https://github.com/langchain-ai/langchain/blob/master/cookbook/rag_fusion.ipynb

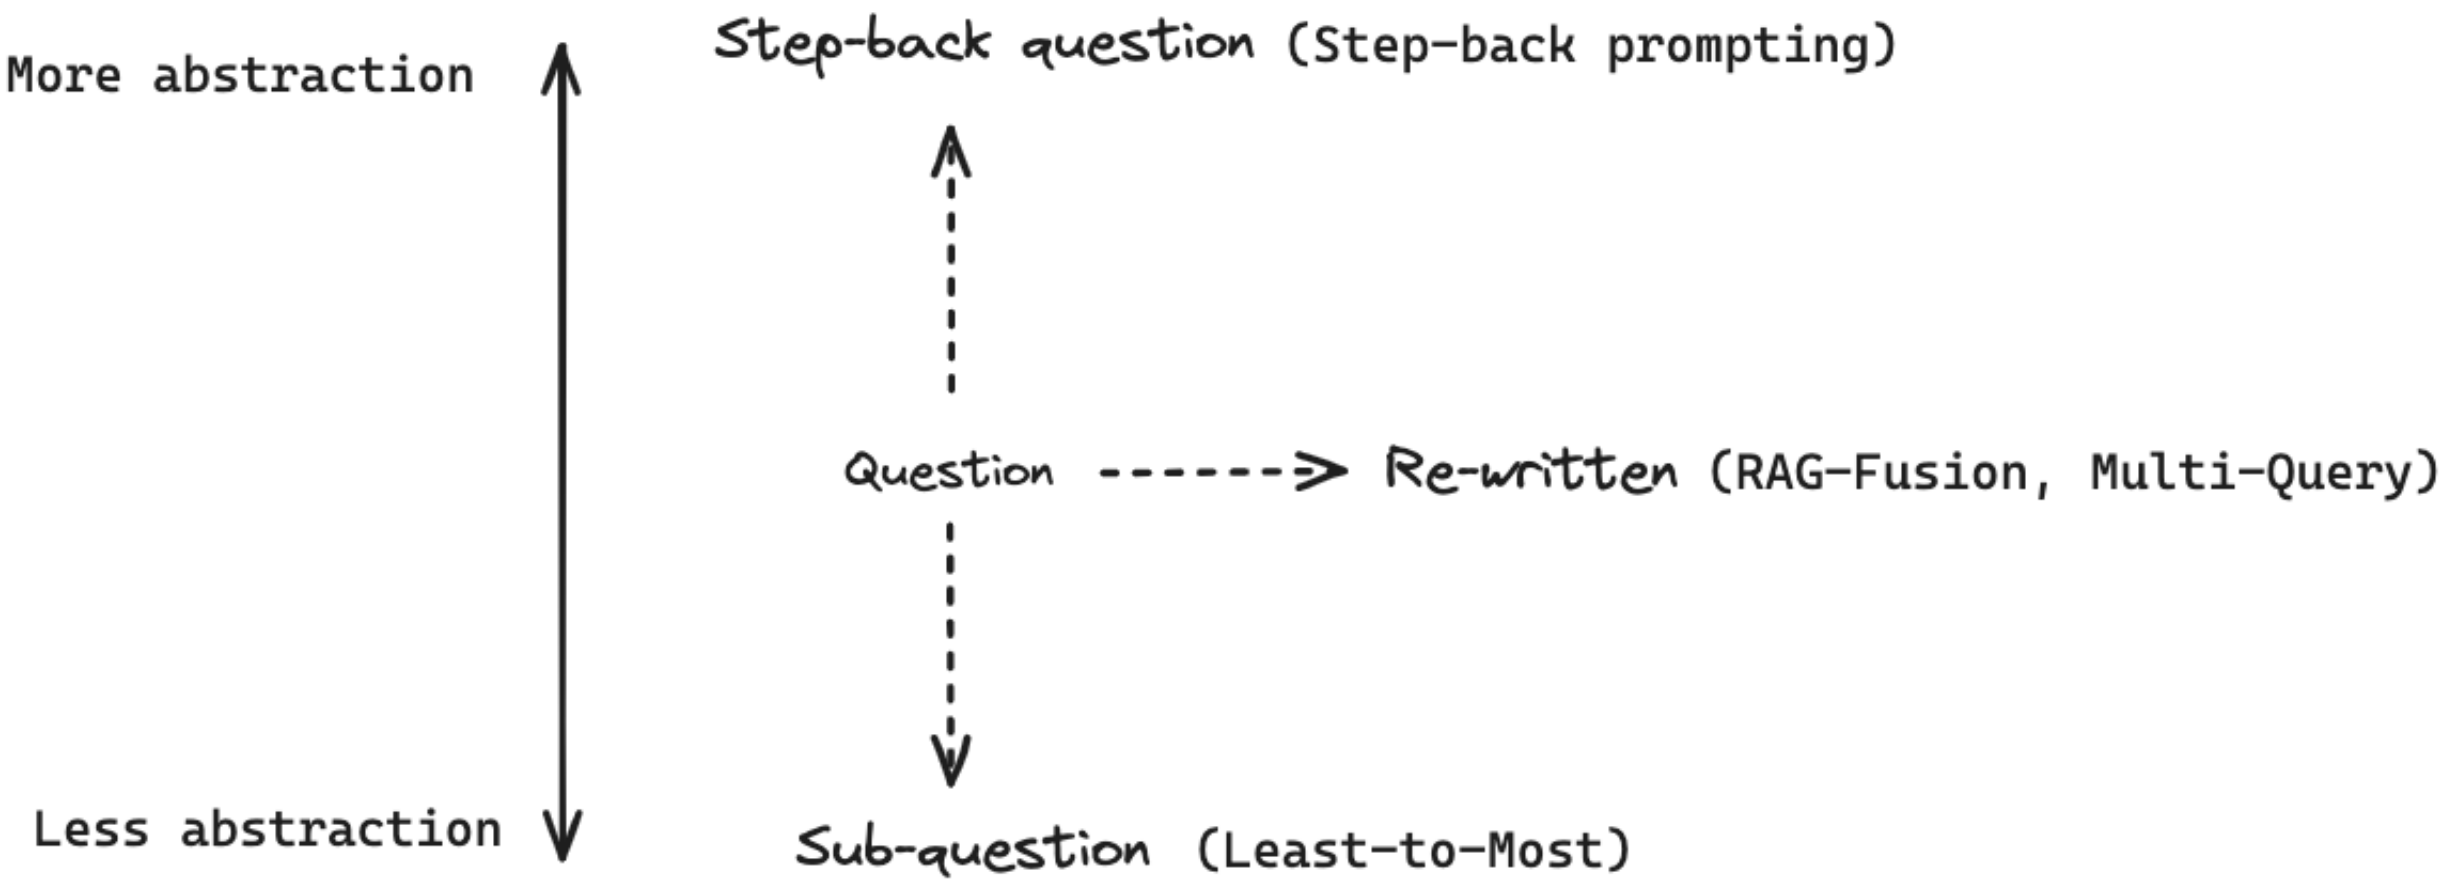
Decomposition:



<https://arxiv.org/pdf/2205.10625.pdf>

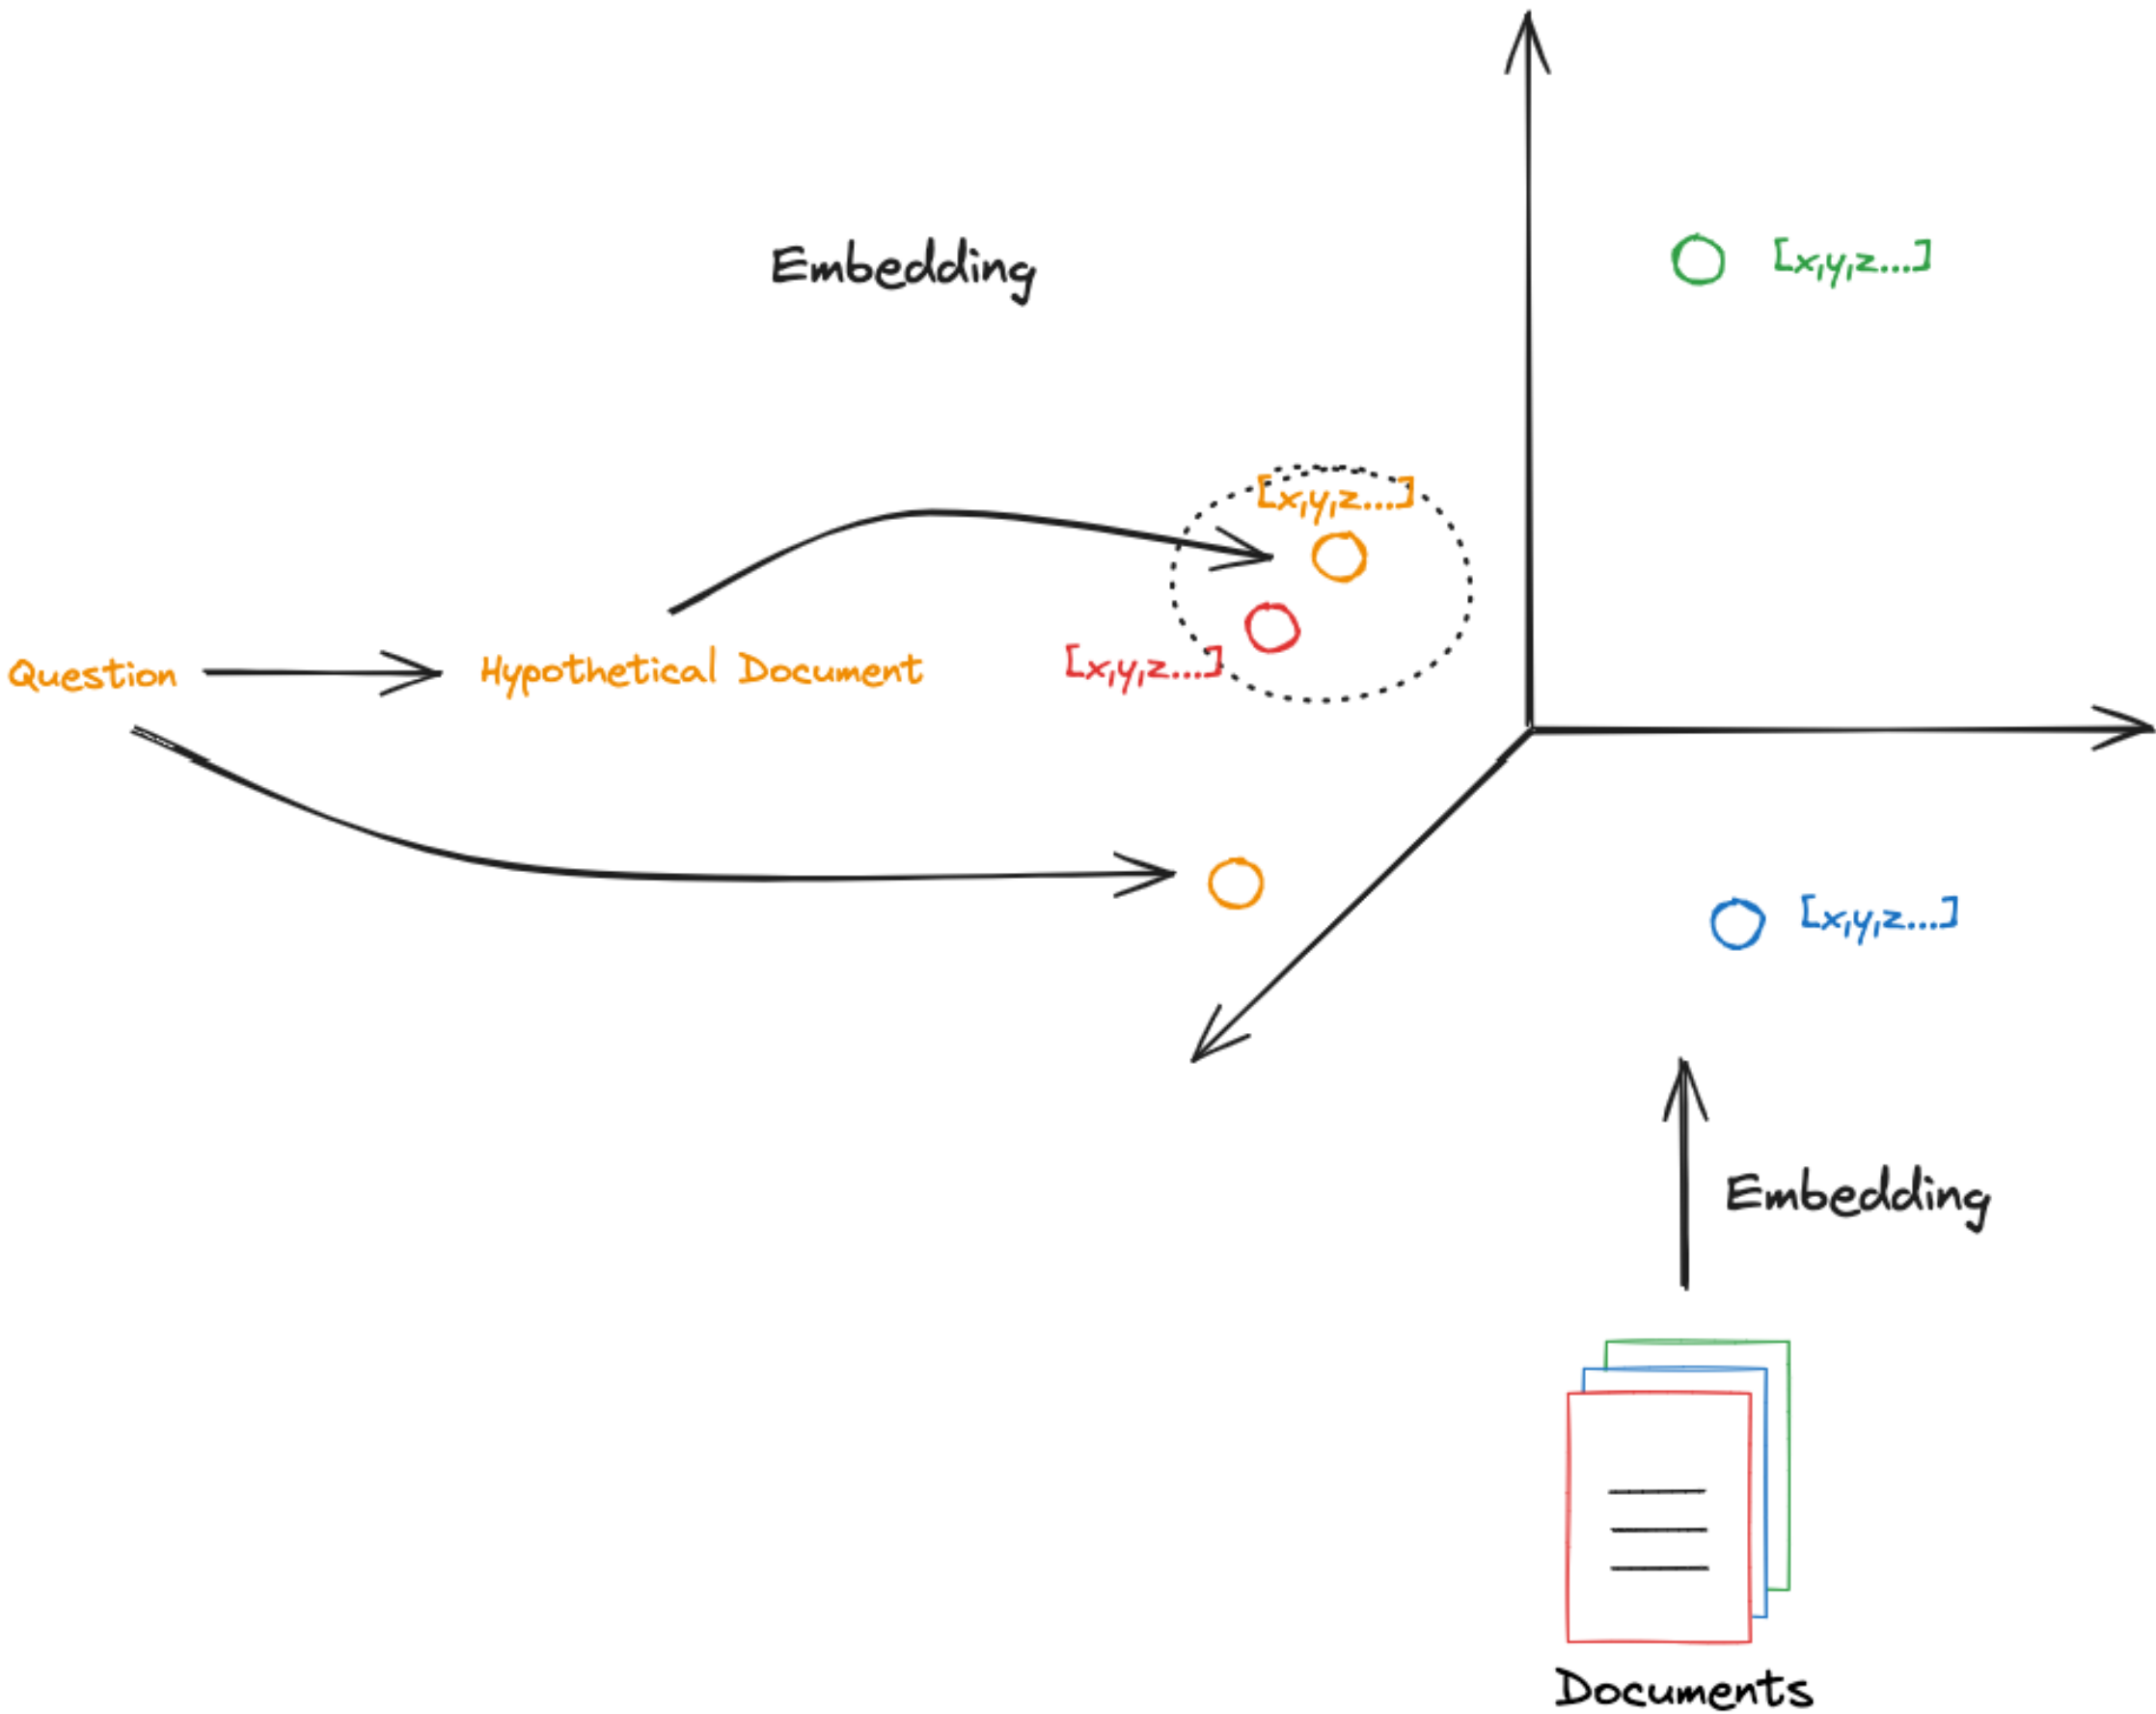
RAG (Retrieval Augmented Generation) Cheatsheet

Step Back:



<https://arxiv.org/pdf/2310.06117.pdf>

HyDE:



<https://arxiv.org/abs/2212.10496>

RAG(Retrieval Augmented Generation)Cheatsheet

Techniques and Tools:

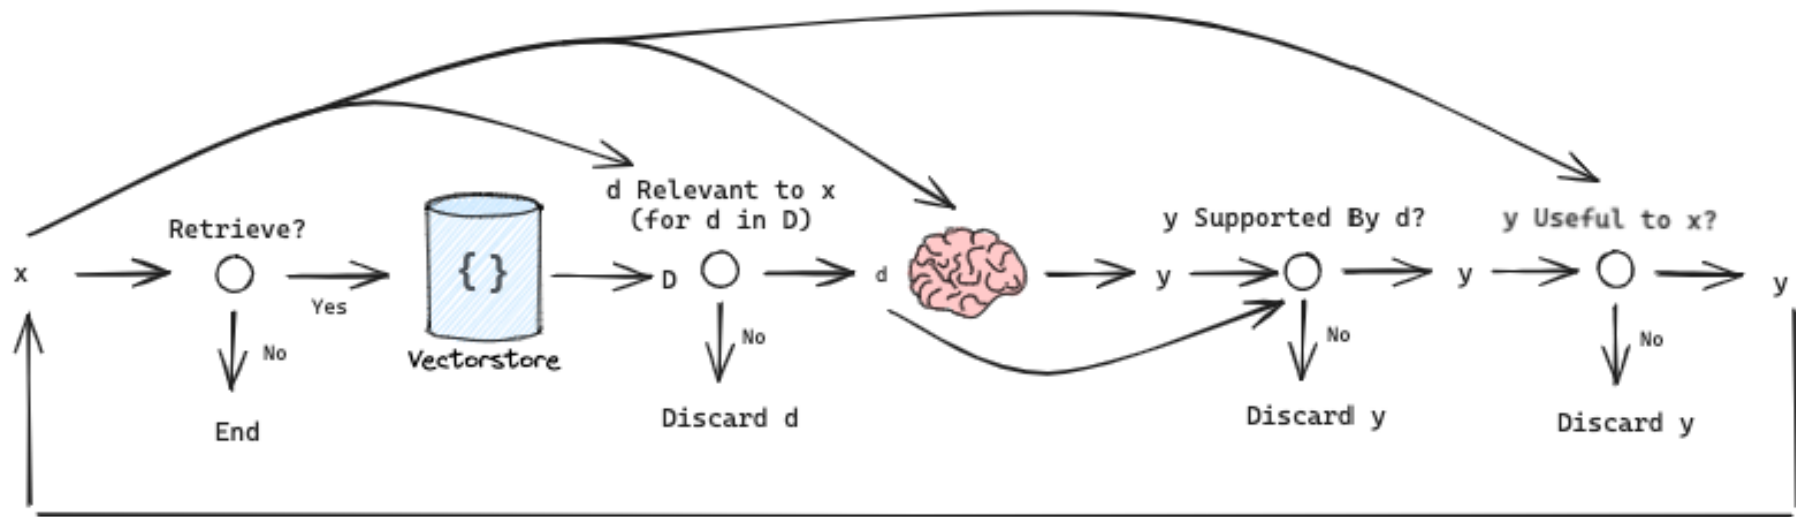
- 1. **Data Ingestion and Querying:**
 - Using tools like LlamaIndex for processing and querying data from various sources into the model's prompt.
- 2. **Chunk Size Optimization:**
 - Adjusting the size of data chunks for efficient processing and retrieval, improving response quality.
- 3. **Metadata Filtering:**
 - Enhancing retrieval by adding structured context to data, utilizing vector database capabilities for more relevant results.
- 4. **Fine-Tuning Embeddings:**
 - Customizing embedding models to better match query context with relevant data, improving precision and recall.
- 5. **Advanced Retrieval Algorithms:**
 - Implementing sophisticated retrieval methods like recursive retrieval and parent-child chunk retrieval to enhance context understanding and response accuracy.

Challenges and Solutions:

- **Missing Data:**
Addressed by expanding the document corpus or integrating external knowledge bases.
- **The issue with Ranking:**
Overcome by using advanced retrieval techniques like rerankers.
- **Consolidation Issues:**
Solved by employing strategies that ensure relevant documents are included in the final context.
- **Formatting Issues:**
Addressed by ensuring the system correctly interprets and responds to format-specific queries.
- **Incorrect Specifics and Incomplete Answers:**
Mitigated by adjusting the detail level of responses to match user queries.
- **Extraction Challenges:**
Overcome by refining the system's ability to accurately extract information from the selected context.

Self-RAG

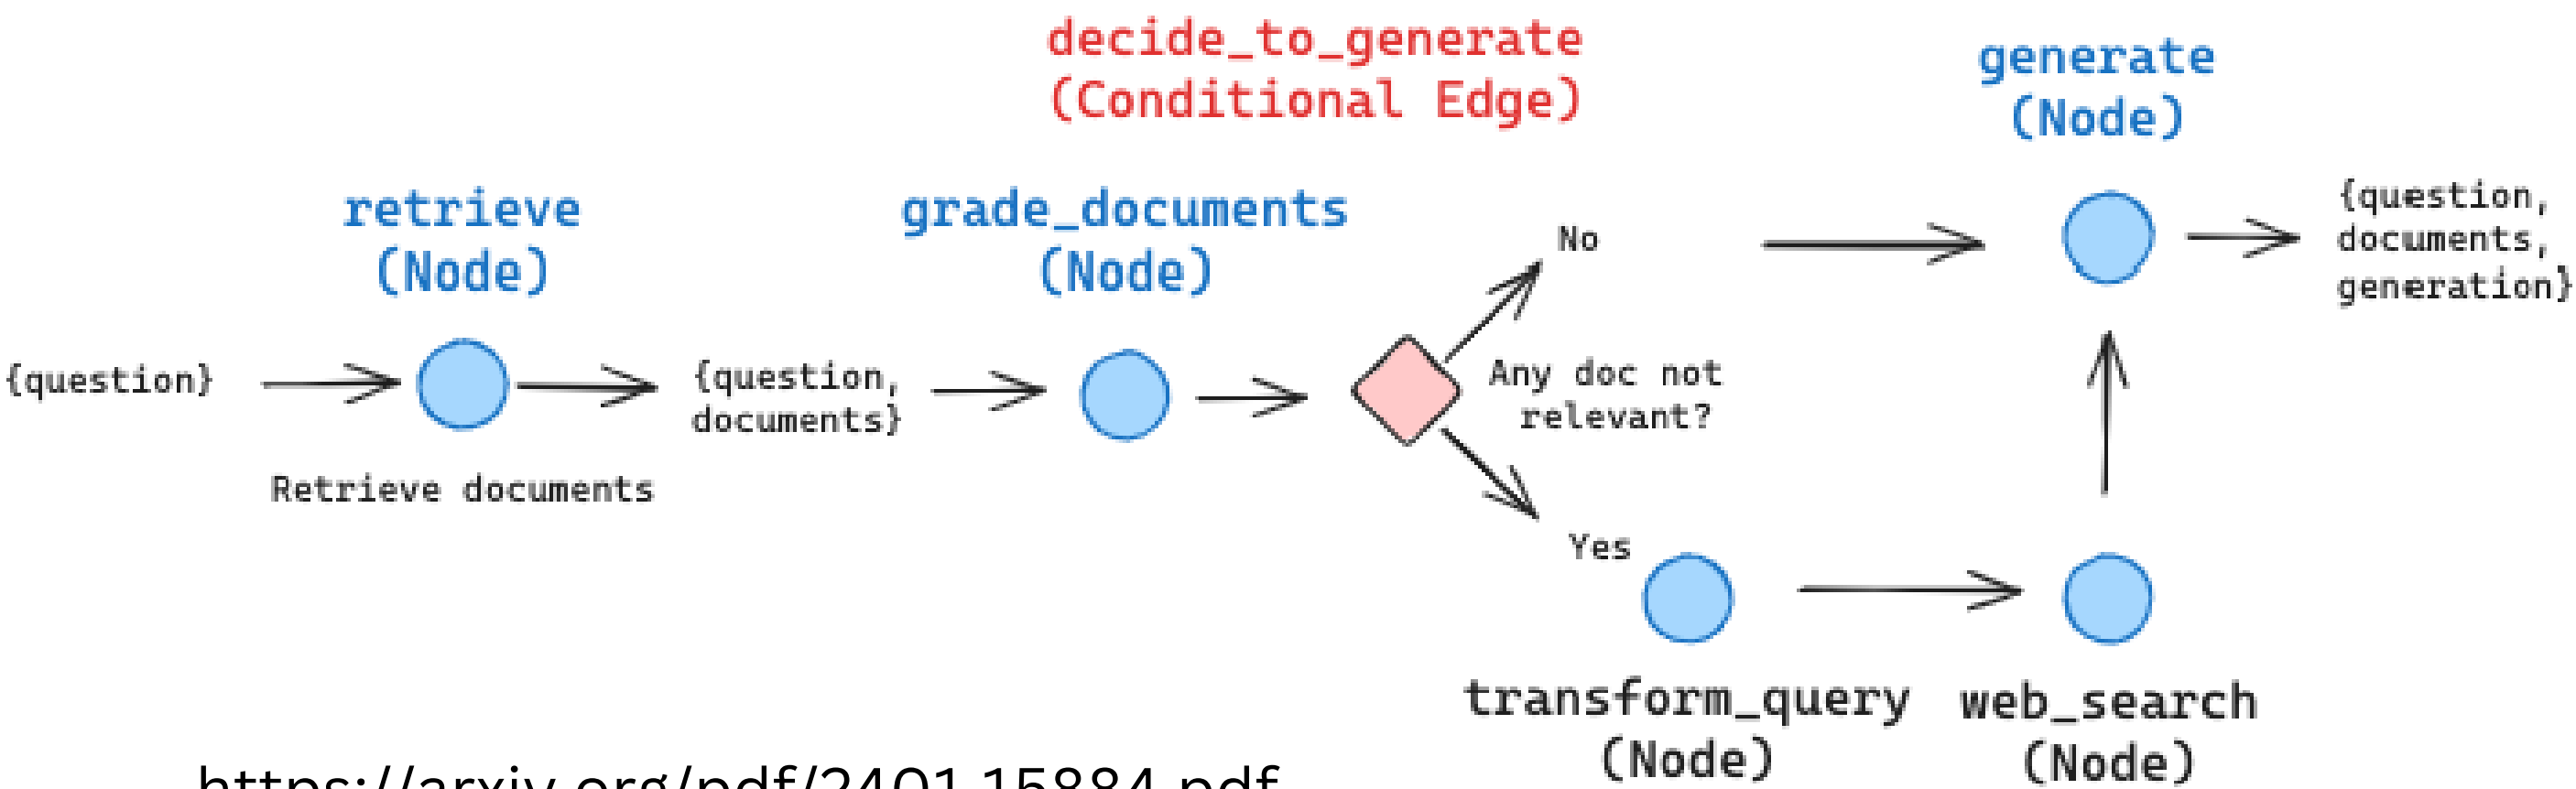
Self-reflection can enhance RAG, enabling correction of poor quality retrieval or generations.



<https://arxiv.org/abs/2310.11511>

Corrective RAG

Corrective-RAG (CRAG) is a recent paper that introduces an interesting approach for self-reflective RAG.



<https://arxiv.org/pdf/2401.15884.pdf>